

Socket Programming

CSEE, Handong Univ.

Jong-won Lee

ljw@handong.edu

Network Programming Interface (API)

□ API

- application callable services
- interfaces and abstractions provided by the system to the application

□ Network Programming at different levels

- send Ethernet, ATM, ... packets
- exchange UDP/TCP packets
- RPC, Xlib, Corba,...

API for TCP/IP

- ❑ TCP/IP does not include an API definition.

- ❑ There are a variety of APIs for use with TCP/IP:
 - Sockets
 - TLI, XTI
 - Winsock
 - MacTCP

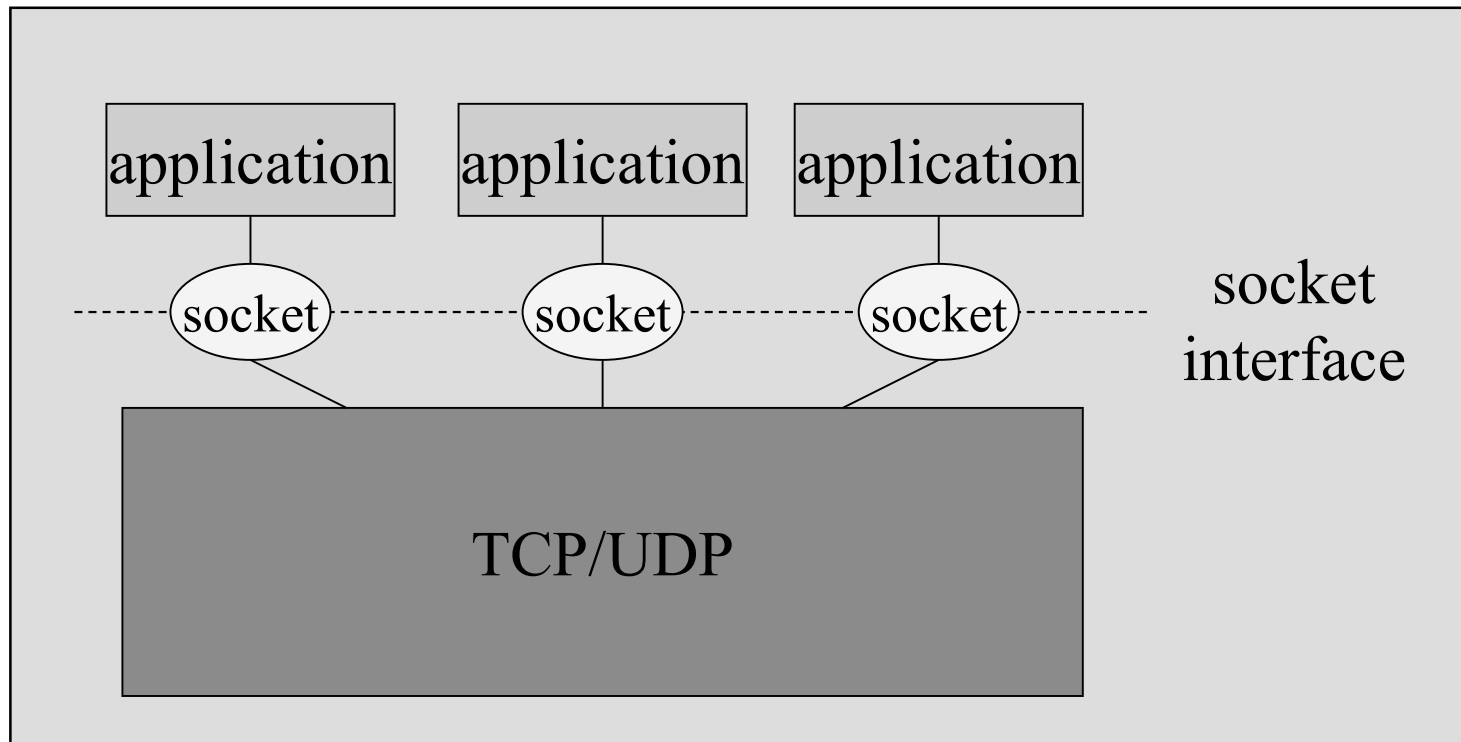
Socket API

- ❑ Introduce in 1981 by BSD 4.1
 - ❑ support for multiple protocol families.
 - ❑ originally only UNIX (=> winsock)
 - ❑ implemented as system calls(BSD) or library(SVR4)
-
- ❑ for TCP/IP, three socket types:
 - stream-oriented: TCP
 - datagram: UDP
 - raw IP: IP, ICMP

Socket


□ socket interface:

- located between application and TCP, UDP and other protocol stacks (common interface)



Socket

□ socket and TCP/UDP relationship



Socket

❑ Each socket is associated with five components.

- Protocol

 - ◆ protocol family and protocol

- source address, source port

- destination address, destination port

❑ Where to define components

- protocol: `socket()`

- source address and port: `bind()`

- destination address and port: `connect()`
`sendto()`

Creating a socket

□ Synopsis

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int socket(int family, int type, int protocol)
```

□ The `socket()` system call returns a socket descriptor (small integer) or a -1 on error.

□ `socket()` allocates resources needed for a communication endpoint

Creating a socket

❑ Synopsis

`int socket(int family, int type, int protocol)`

❑ `family` specifies the protocol family

- `PF_INET` : Internet protocol (TCP/IP)

- `PF_INET6` : IPv6

- `PF_LOCAL` : for local communication

- `PF_UNIX` : UNIX system internal protocol

❑ `type` specifies the type of service

- `SOCK_STREAM` :

- `SOCK_DGRAM` :

- `SOCK_RAW` : raw IP

❑ `protocol` specifies the specific protocol

- (usually 0 which means the default).

Creating a socket

Example

```
#include <stdio.h>
#include <sys/types.h>
#include <sys/socket.h>

int sockfd;

if ((sockfd = socket(PF_INET, SOCK_STREAM, 0)) < 0) {
    /* print "socket error + the error message */
    perror("socket error"); exit(1);
}
```

Print error : perror()

```
#include <stdio.h>
```

```
void perror(const char *s)
```

Bind the local address

□ Synopsis

```
#include <sys/types.h>
#include <sys/socket.h>
```

```
int bind(int sockfd, struct sockaddr *addr, int addr_len)
```

□ The `bind()` system call is used to assign an address to an existing socket.

□ `Bind()` returns 0 if successful or -1 on error.

Address Structures

❑ Defined in <netinet/in.h>

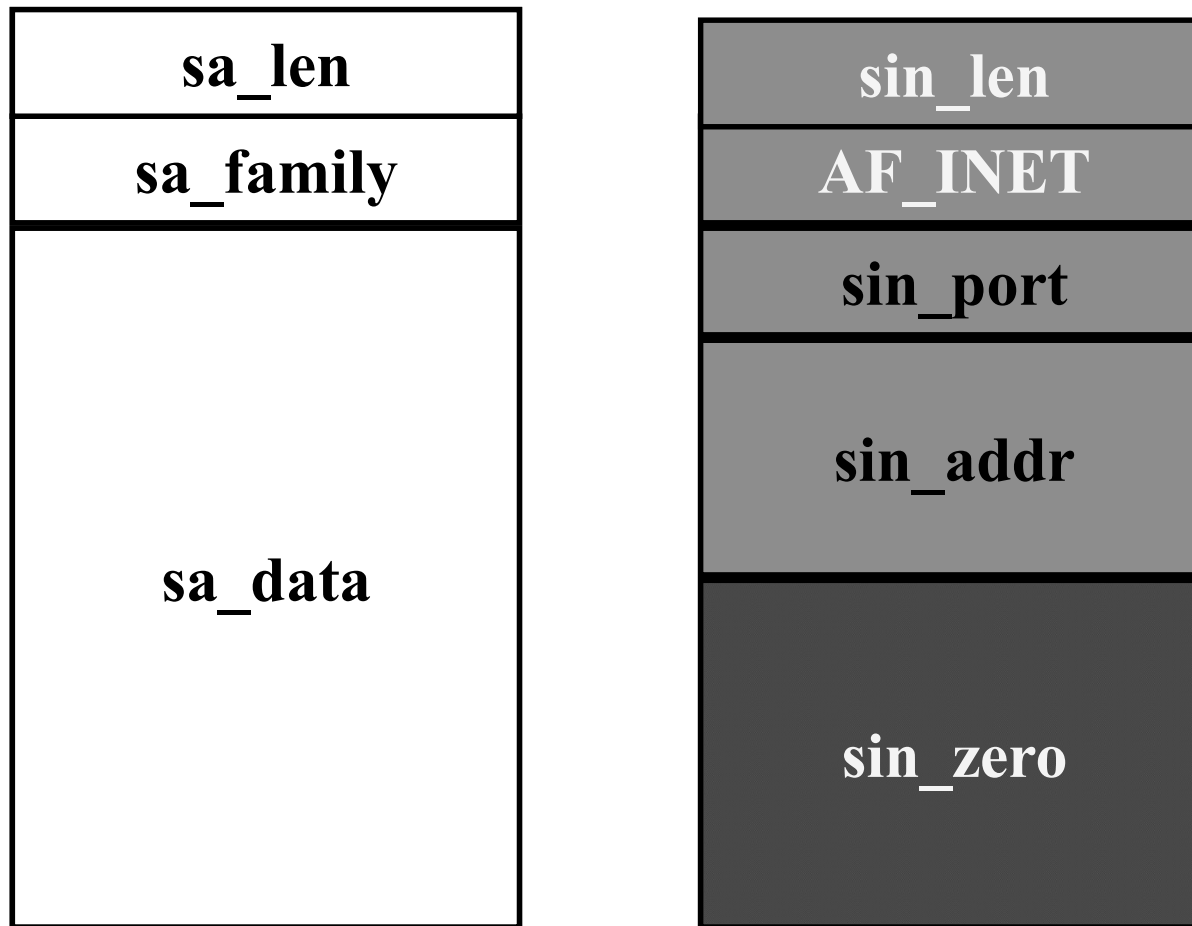
```
struct sockaddr {  
    u_char      sa_len;           /* length : used in kernel */  
    u_short     sa_family;       /* address family */  
    char        sa_data[14];     /* address */  
}
```

```
struct sockaddr_in {  
    u_char      sin_len;         /* length */  
    u_short     sin_family;      /* AF_INET */  
    u_short     sin_port;       /* port number */  
    struct in_addr sin_addr;     /* IP address */  
    char        sin_zero[8];    /* unused */  
}
```

```
struct in_addr {  
    u_long      s_addr; /* 32 bit IP address */  
}
```

Address Structures

□ sockaddr vs. sockaddr_in



Bind() Example

```
#include <stdio.h>
#include <string.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa/inet.h>

#define MYPORT 50000

int sockfd;
struct sockaddr_in my_addr;

if ((sockfd = socket(PF_INET, SOCK_STREAM, 0)) < 0) {
    perror("socket error"); exit (1); }
memset(&my_addr, 0, sizeof(my_addr));
/* bzero( (char *)&my_addr, sizeof(my_addr)); */
my_addr.sin_family = AF_INET;
my_addr.sin_port = htons(MYPORT);
my_addr.sin_addr.s_addr = htonl(INADDR_ANY);
if (bind(sockfd, (struct sockaddr *)&my_addr, sizeof(my_addr)) < 0) {
    perror("bind error" ); exit(1);}
```

Byte Order Conversion

❑ Byte ordering

- little endian: least significant byte first (Intel)
- big endian: most significant byte first (Motorola, Sun)

❑ *network byte order* = big endian

❑ Synopsis

```
#include <sys/types.h>
#include <netinet/in.h>

/* byte swapping host byte order <-> network */
/* for long and short integer */
u_long    htonl(u_long hostlong);

u_long ntohl(u_long netlong);

u_short htons(u_short hostshort);

u_short ntohs(u_short netshort)
```

Name-to-Address Conversion

Synopsis

```
#include <sys/socket.h>  
#include <netdb.h>
```

```
/* return host information by taking host name */  
struct hostent *gethostbyname(const char *name);
```

```
/* return host information by taking network byte order  
address */  
struct hostent *gethostbyaddr(const char *addr, int  
length, int type);
```


Looking up A Domain Name

Defined in <netdb.h>

```
#define h_addr haddr_list[0]
struct hostent {
    char    *h_name;        /* host name */
    char    **h_alias;       /* list of alternate names */
    int     h_addrtype;     /* type of address = 2 (=AF_INET) */
    int     h_length;       /* address length = 4 for IPv4 */
    char    **h_addr_list;; /* address list */
};
```

Gethostbyname()

Example

```
struct hostent *phost;
struct in_addr **addr_list;

if ((phost = gethostbyname("www.handong.edu")) == NULL) {
    perror("gethostbyname");
    return 1;
}
// print information about this host:
printf("Official name is: %s\n", phost->h_name);
printf("  IP addresses: ");
addr_list = (struct in_addr **)phost->h_addr_list;
for(i = 0; addr_list[i] != NULL; i++) {
    printf("%s ", inet_ntoa(*addr_list[i]));
}
printf("\n");
```

Gethostbyaddr()

Example

```
struct hostent *phost;  
struct in_addr addr;  
  
inet_aton("203.252.97.12", &addr);  
phost = gethostbyaddr(&addr, sizeof(addr), AF_INET);  
printf("Host name: %s\n", phost->h_name);
```

Looking up A Well-known Port by Name

Defined in <netdb.h>

```
struct servent {  
    char    *s_name;        /* official service name */  
    char    **s_alias;      /* list of alternate names */  
    int     s_port;         /* port for this service */  
    int     *s_proto;       /* protocol to use */  
};
```

```
struct servent *pService;
```

```
if (pService = getservbyname("smtp", "tcp")) {  
    // port number is now in pService->s_port;  
} else {  
    /* error */  
}
```

New Functions: name-to-addr conversion

getaddrinfo(): converts human-readable text strings representing hostname into a dynamically allocated linked list of struct addrinfo structures.

getnameinfo(): looks up the host name and service name information for a given struct sockaddr.

Synopsis

```
#include <sys/socket.h>
```

```
#include <netdb.h>
```

```
int getaddrinfo(const char *name, const char *servname,  
               const struct addrinfo *hints, struct addrinfo **res);
```

```
int getnameinfo(const struct sockaddr *sa, socklen_t salen,  
               char *host, size_t hostlen,  
               char *serv, size_t servlen, int flags)
```

New Functions: name-to-addr conversion

Struct addrinfo

```
struct addrinfo {  
    int    ai_flags;           // AI_PASSIVE, AI_CANONNAME, ...  
    int    ai_family;         // AF_XXX  
    int    ai_socktype;       // SOCK_XXX  
    int    ai_protocol;       // 0 (auto) or IPPROTO_TCP, IPPROTO_UDP  
    socklen_t ai_addrlen;     // length of ai_addr  
    char *ai_canonname;       // canonical name for nodename  
    struct sockaddr *ai_addr; // binary address  
    struct addrinfo *ai_next; // next structure in linked list  
};
```

Ge

E

```
int sockfd;
struct addrinfo hints, *servinfo, *p;
int rv;
memset(&hints, 0, sizeof hints);
hints.ai_family = AF_UNSPEC;    // use AF_INET6 to force IPv6
hints.ai_socktype = SOCK_STREAM;

if ((rv = getaddrinfo("www.example.com", "http", &hints, &servinfo)) != 0) {
    fprintf(stderr, "getaddrinfo: %s\n", gai_strerror(rv));
    exit(1);
}

for(p = servinfo; p != NULL; p = p->ai_next) {
    if ((sockfd = socket(p->ai_family, p->ai_socktype, p->ai_protocol)) == -1) {
        perror("socket");
        continue;
    }
    if (connect(sockfd, p->ai_addr, p->ai_addrlen) == -1) {
        close(sockfd);
        perror("connect");
        continue;
    }
    break; // if we get here, we must have connected successfully
}
...

freeaddrinfo(servinfo); // all done with this structure
```

Jong-w

Getnameinfo()

Example

```
struct sockaddr_in sa;  
char host[1024];  
char service[20];  
  
//pretend sa is full of good information about the host and port...  
  
getnameinfo(&sa, sizeof sa, host, sizeof (host), service, sizeof (service), 0);  
  
printf("  host: %s\n", host);  
printf("service: %s\n", service);
```


IP Address Manipulation

Dotted decimal vs. IP address

Synopsis

```
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
```

```
/* IP address to dotted decimal */
char *inet_ntoa(struct in_addr address);
```

```
/* dotted decimal to IP address */
u_long inet_addr(char *dottedAddress);
    // return -1 if error
Int inet_aton(const char *cp, struct in_addr *inp);
    // return 0 if error
```

IP Address Manipulation

Synopsis

```
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#include <arpa.h>
```

```
struct sockaddr_in  my_addr;
```

```
memset(&my_addr, 0, sizeof(my_addr));
my_addr.sin_family = AF_INET;
my_addr.sin_port = htons(MYPORT);
inet_aton("10.12.110.57", &(my_addr.sin_addr));
/* my_addr.sin_addr.s_addr = inet_addr("10.12.110.57"); */
printf("%s", inet_ntoa(my_addr.sin_addr));
```

Address Format Conversion Summary

Domain name : IP address (binary) : Dotted decimal

gethostbyname()



gethostbyaddr()



inet_ntoa()



inet_addr()



inet_aton()

IP Address Manipulation

Synopsis

```
#include <sys/socket.h>
#include <arpa/inet.h>
```

```
const char *inet_ntop(int af, const void *src,
                      char *dst, socklen_t size);
// return value: null if error
```

```
int inet_pton(int af, const char *src, void *dst)
//return value: -1 if error,
//              0 for invalid IP address
```

Example

```
struct sockaddr_in sa;  
char str[INET_ADDRSTRLEN];  
  
// store this IP address in sa:  
inet_pton(AF_INET, "192.0.2.33", &(sa.sin_addr));  
  
// now get it back and print it  
inet_ntop(AF_INET, &(sa.sin_addr), str, INET_ADDRSTRLEN);  
printf("%s\n", str); // prints "192.0.2.33"
```

Address Format Conversion Summary

Domain name : IP address (binary) : Dotted decimal

`getaddrinfo()`



`inet_ntop()`



`getnameinfo()`



`inet_pton()`



System Call : connect()

Initiate a connection on a socket (client only)

- Returns 0 on success; -1 on failure
- For a TCP socket, it establishes a connection to the server
- For a UDP socket, it simply stores the server's address so that the client can use a socket description

Synopsis

```
#include <sys/types.h>
#include <sys/socket.h>
```

```
int connect(int sd, struct sockaddr *addr, int addr_len);
/* addr = server address */
```

Connect() Example

```
#include <string.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>

#define DEST_IP "10.12.110.57"
#define MYPORT 23

int sockfd;
struct sockaddr_in dest_addr;

if ((sockfd = socket(PF_INET, SOCK_STREAM, 0)) < 0) {
    /* error */ }
memset(&dest_addr, 0, sizeof(dest_addr));
dest_addr.sin_family = AF_INET;
dest_addr.sin_port = htons(DEST_PORT);
dest_addr.sin_addr.s_addr = inet_addr(DEST_IP);
If (connect(sockfd, (struct sockaddr *)&dest_addr, sizeof(dest_addr)) != 0 ) {
    close(sockfd);
    return -1;
}
```


System Call : listen()

**Tell OS to receive and queue SYN packets
Specify backlog size for the pending connection
requests on a socket (TCP server only)**

○Returns 0 on success; -1 on failure

Synopsis

#include <sys/types.h>

#include <sys/socket.h>

int listen(int sd, int backlog);

Parameter

○**backlog** : specifies the max. number of connection requests that system can queue while it waits for the server to accept them (usually 5)

System Call : accept()

Accept a connection on a socket

- TCP server only
- Returns a new socket descriptor (>0) on success;
-1 on failure
- Block until a connection request arrives

Synopsis

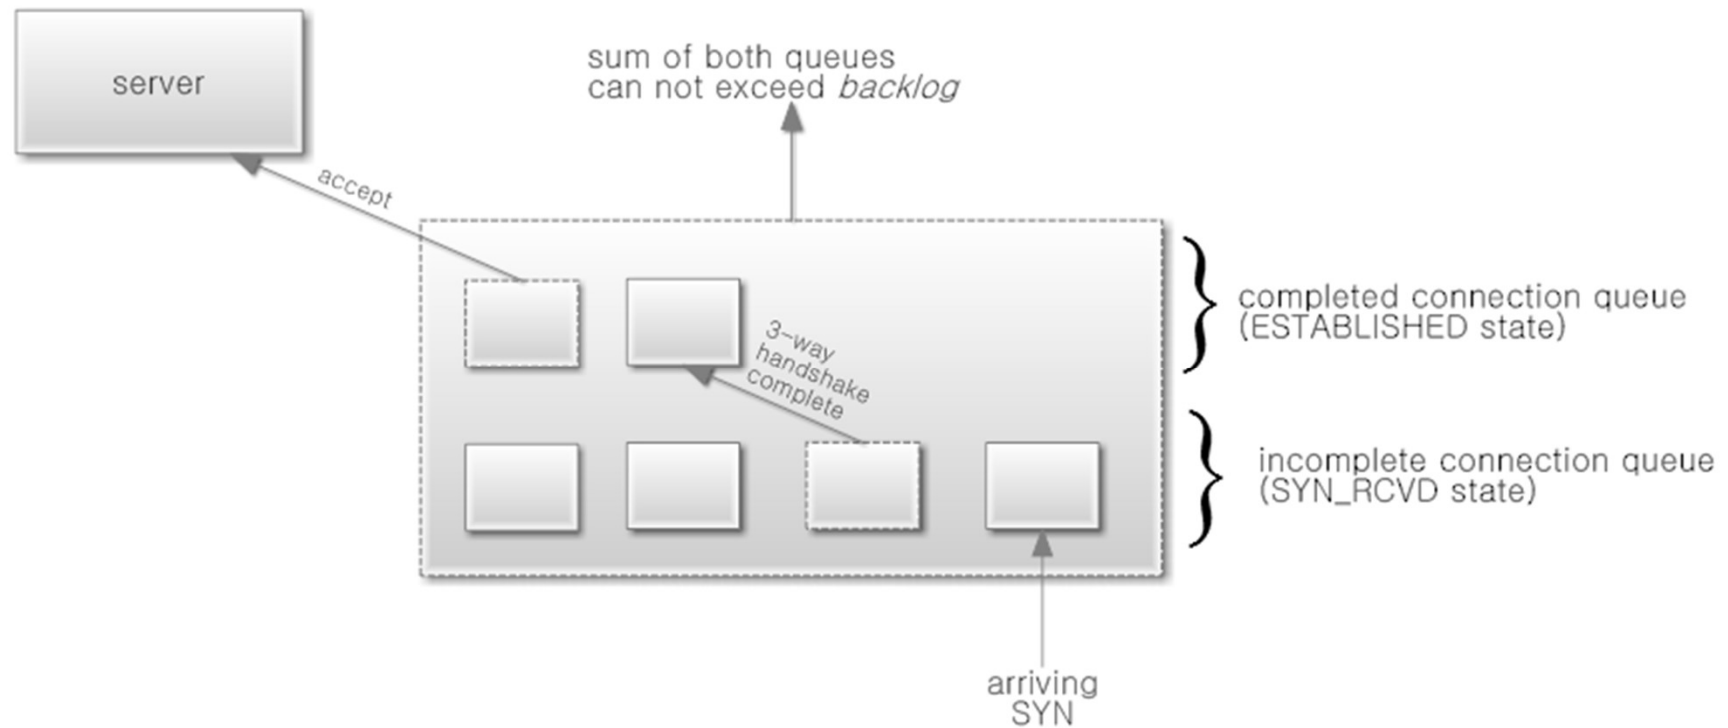
```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int accept(int sd, struct sockaddr *addr, int *addrlen);
```

Parameters

- addr: a pointer to an address structures to be filled in
- addrlen: a pointer to an integer that should be set to
`sizeof(struct sockaddr_in)`



Example of TCP Server

```
#include <string.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>

#define MYPORT 50000
main()
{
    int sockfd, new_fd;
    struct sockaddr_in my_addr, client_addr;
    int sin_size;

    sockfd = socket(PF_INET, SOCK_STREAM, 0);
    memset(&my_addr, 0, sizeof(my_addr));
    my_addr.sin_family = AF_INET;
    my_addr.sin_port = htons(MYPORT);
    my_addr.sin_addr.s_addr = htonl(INADDR_ANY);
    bind(sockfd, (struct sockaddr *)&my_addr, sizeof(my_addr));
    listen(sockfd, 5);
    sin_size = sizeof(client_addr);
    new_fd = accept(sockfd, (struct sockaddr *)&client_addr, &sin_size);
```

Sending and Receiving data

Function calls

- **read(), write() : any device, single buffering**
- **readv(), writev() : any device, from/to several buffers**
- **recv(), send() : socket device, single buffer**
- **recvfrom(), sendto() : socket device, single buffer, specifying peer**

System calls : read() and write()

read()

- `int read(int fd, char *buf, int buflen)`
- Returns the number of bytes received on success;
-1 on failure
 - ◆ When 0 is returned, it means that the remote side has closed the connection.
- Block until data received

write()

- `int write(int fd, char *buf, int buflen)`
- Returns the number of bytes transmitted on success;
-1 on failure

System calls : send()

send()

`#include <sys/types.h>`

`#include <sys/socket.h>`

`int send(int sockfd, const void *msg, int len, int flag)`

- Returns the number of bytes transmitted on success; -1 on failure
- msg: the pointer to the data to send
- flag: 0 for normal data

○Example:

```
char *msg = "Hi, Beej!";  
int len, bytes_sent;  
len = strlen(msg);  
byte_sent = send(sockfd, msg, len, 0);
```

System calls : recv()

synopsis

#include <sys/types.h>

#include <sys/socket.h>

int recv(int sockfd, void *buf, int len, unsigned int flag)

○ **Returns the number of bytes received on success; -1 on failure**

◆ **When 0 is returned, it means that the remote side has closed the connection.**

○ **Block until data received**

○ **buf: the buffer to read the information into**

○ **len: the max. length of the buffer**

○ **flag: 0 for regular data**

System Call : sendto()

Synopsis

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int sendto(int sockfd, const void *msg, int len, int flag,  
           struct sockaddr *dstaddr, int addrlen)
```

○Returns the number of bytes sent on success; -1 on failure

Parameters

- Transmit the data in *msg* upto *len* bytes
- flags** : to control transmission behavior (0 for normal)
- The **dstaddr** includes the information of destination address

System Call : recvfrom()

synopsis

#include <sys/types.h>

#include <sys/socket.h>

**int recvfrom(int sockfd, void *buf, int len, int flags,
 struct sockaddr *srcaddr, int *addrlen)**

- **Returns the number of bytes received on success;
-1 on failure**
- **Blocks until data received (normal operation)**

Parameters

- **Receives up to *len* bytes into *buf***
- **flags : to control transmission behavior (0 for normal)**
- **The *srcaddr* is filled in with the address of the sender.**

System Call : close()

Synopsis

#include <unistd.h>

int close(int sockfd)

- **Prevent any more reads and writes to the socket.**
 - ◆ **if the remote side calls `recv()`, it will return 0.**
 - ◆ **if the remote side calls `send()`, it'll receive a signal `SIGPIPE` and `send()` will return -1 and `errno` will be set to `EPIPE`.**

System Call : shutdown()

Synopsis

#include <sys/socket.h>

Int shutdown(int sockfd, int how)

- **Returns 0 on success, and -1 on error**
- **The operation depends on the value of how**
 - ◆ **0 (SHUT_RD): further receives are disallowed**
 - ◆ **1 (SHUT_WR): further sends are disallowed**
 - ◆ **2 (SHUT_RDWR): further sends and receives are disallowed**
- **Note that shutdown() does not actually free up the socket descriptor. To free the descriptor, use close().**

Half-close: shutdown()

Can notify the end of file transmission.

